

Python for HPC

Andrea Zonca - SDSC

Jupyter Notebook

Data exploration in your browser

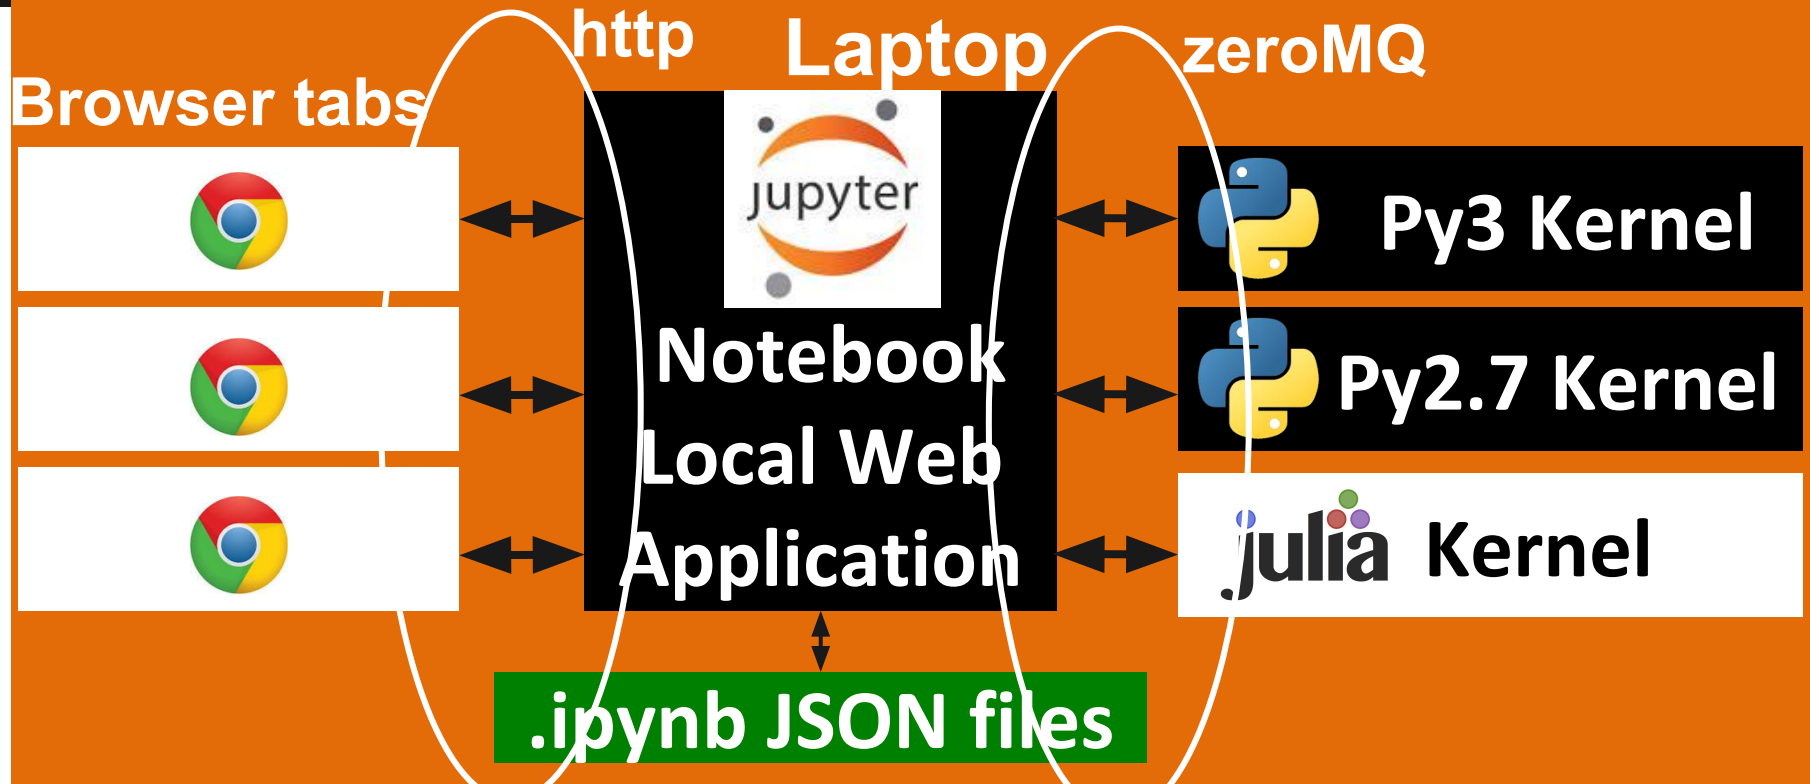
What is the notebook?

- Browser based interactive console
- Supports multiple sessions in browser tabs
- Each session has a Kernel executing computation
- Saved in JSON format

Notebooks on Nature

<http://www.nature.com/news/interactive-notebooks-sharing-the-code-1.16261>

Jupyter notebook local



Jupyter notebook remote

Laptop



https +
password

Server

Jupyter
Notebook
Web
Application



Py3 Kernel



Py2.7 Kernel



Kernel

.ipynb JSON files

Clone workshop repository

ssh into comet with training account

git clone URL

Find url at the address bit.ly/cometucsb

More secure setup

<http://zonca.github.io/2015/09/ipython-jupyter-notebook-sdsc-comet.html>

IPython notebook demo

- Python code
- Formatted text
- Equations
- Plots
- Cells execution, cells order
- Clear output

Why the notebook?

- Literate programming: code and explanation together
- Reproducible science: document easily every step
- Easy to share computations: send one single notebook instead of scripts/plots/.doc

ipynb documents

- JSON format
- includes plots in binary format
- easy to convert to .html/.pdf for sharing
- <http://nbviewer.ipython.org>
- Recently rendered automatically on Github

HPC: interactive notebooks

- Analyze large amount of data
- In-situ visualization
- Centralized Python stack
- Check long-running computations
- Prepare and submit batch jobs

Notebooks as scripts

- Run notebooks not-interactively:

```
jupyter nbconvert --execute --to notebook  
input_notebook.ipynb --output  
executed_notebook.ipynb
```

Dask

Multithreading and Distributed computing

Numba

Run code on GPU with Python

JIT compiler for Python

- based on LLVM (compiler infrastructure behind clang, Apple's C++ compiler)
- turns Python code into machine code
- on-the-fly

Numba

export

NUMBAPRO_NVVM=/usr/local/cuda-7.0/nvvm/lib64/libnvvm.so

export

NUMBAPRO_LIBDEVICE=/usr/local/cuda-7.0/nvvm/libdevice/

Interactive GPU node

```
salloc --nodes=1 --tasks-per-node=24  
--partition=gpu -t 01:00:00
```

```
from numba import jit
from numpy import arange
```

```
# jit decorator tells Numba to compile this function.
```

```
# The argument types will be inferred by Numba when function is called.
```

```
@jit
```

```
def sum2d(arr):
```

```
    M, N = arr.shape
```

```
    result = 0.0
```

```
    for i in range(M):
```

```
        for j in range(N):
```

```
            result += arr[i,j]
```

```
    return result
```

```
a = arange(9).reshape(3,3)
```

```
print(sum2d(a))
```

Numba CPU

run with `%timeit`

increase size of matrix to see performance improvements

Numba GPU

```
from numba import cuda

@cuda.jit
def matmul(A, B, C):
    """Perform square matrix multiplication of  $C = A * B$ 
    """
    i, j = cuda.grid(2)
    if i < C.shape[0] and j < C.shape[1]:
        tmp = 0.
        for k in range(A.shape[1]):
            tmp += A[i, k] * B[k, j]
        C[i, j] = tmp

import numpy as np
shape = (5,5)
a = np.ones(shape)
b = np.ones(shape) * 4
c = np.zeros(shape)
matmul[1,(16,16)](a,b,c)
print(c)
```

Hands-on

- create a loop that runs `matmul` with different matrix sizes
- compare with `np.dot`
- range from `20x20` to `10000x10000`
- plot timing

Advanced CUDA

Tiled matrix multiplication to exploit GPU fast local memory:

<http://numba.pydata.org/numba-doc/0.27.0/cuda/examples.html>